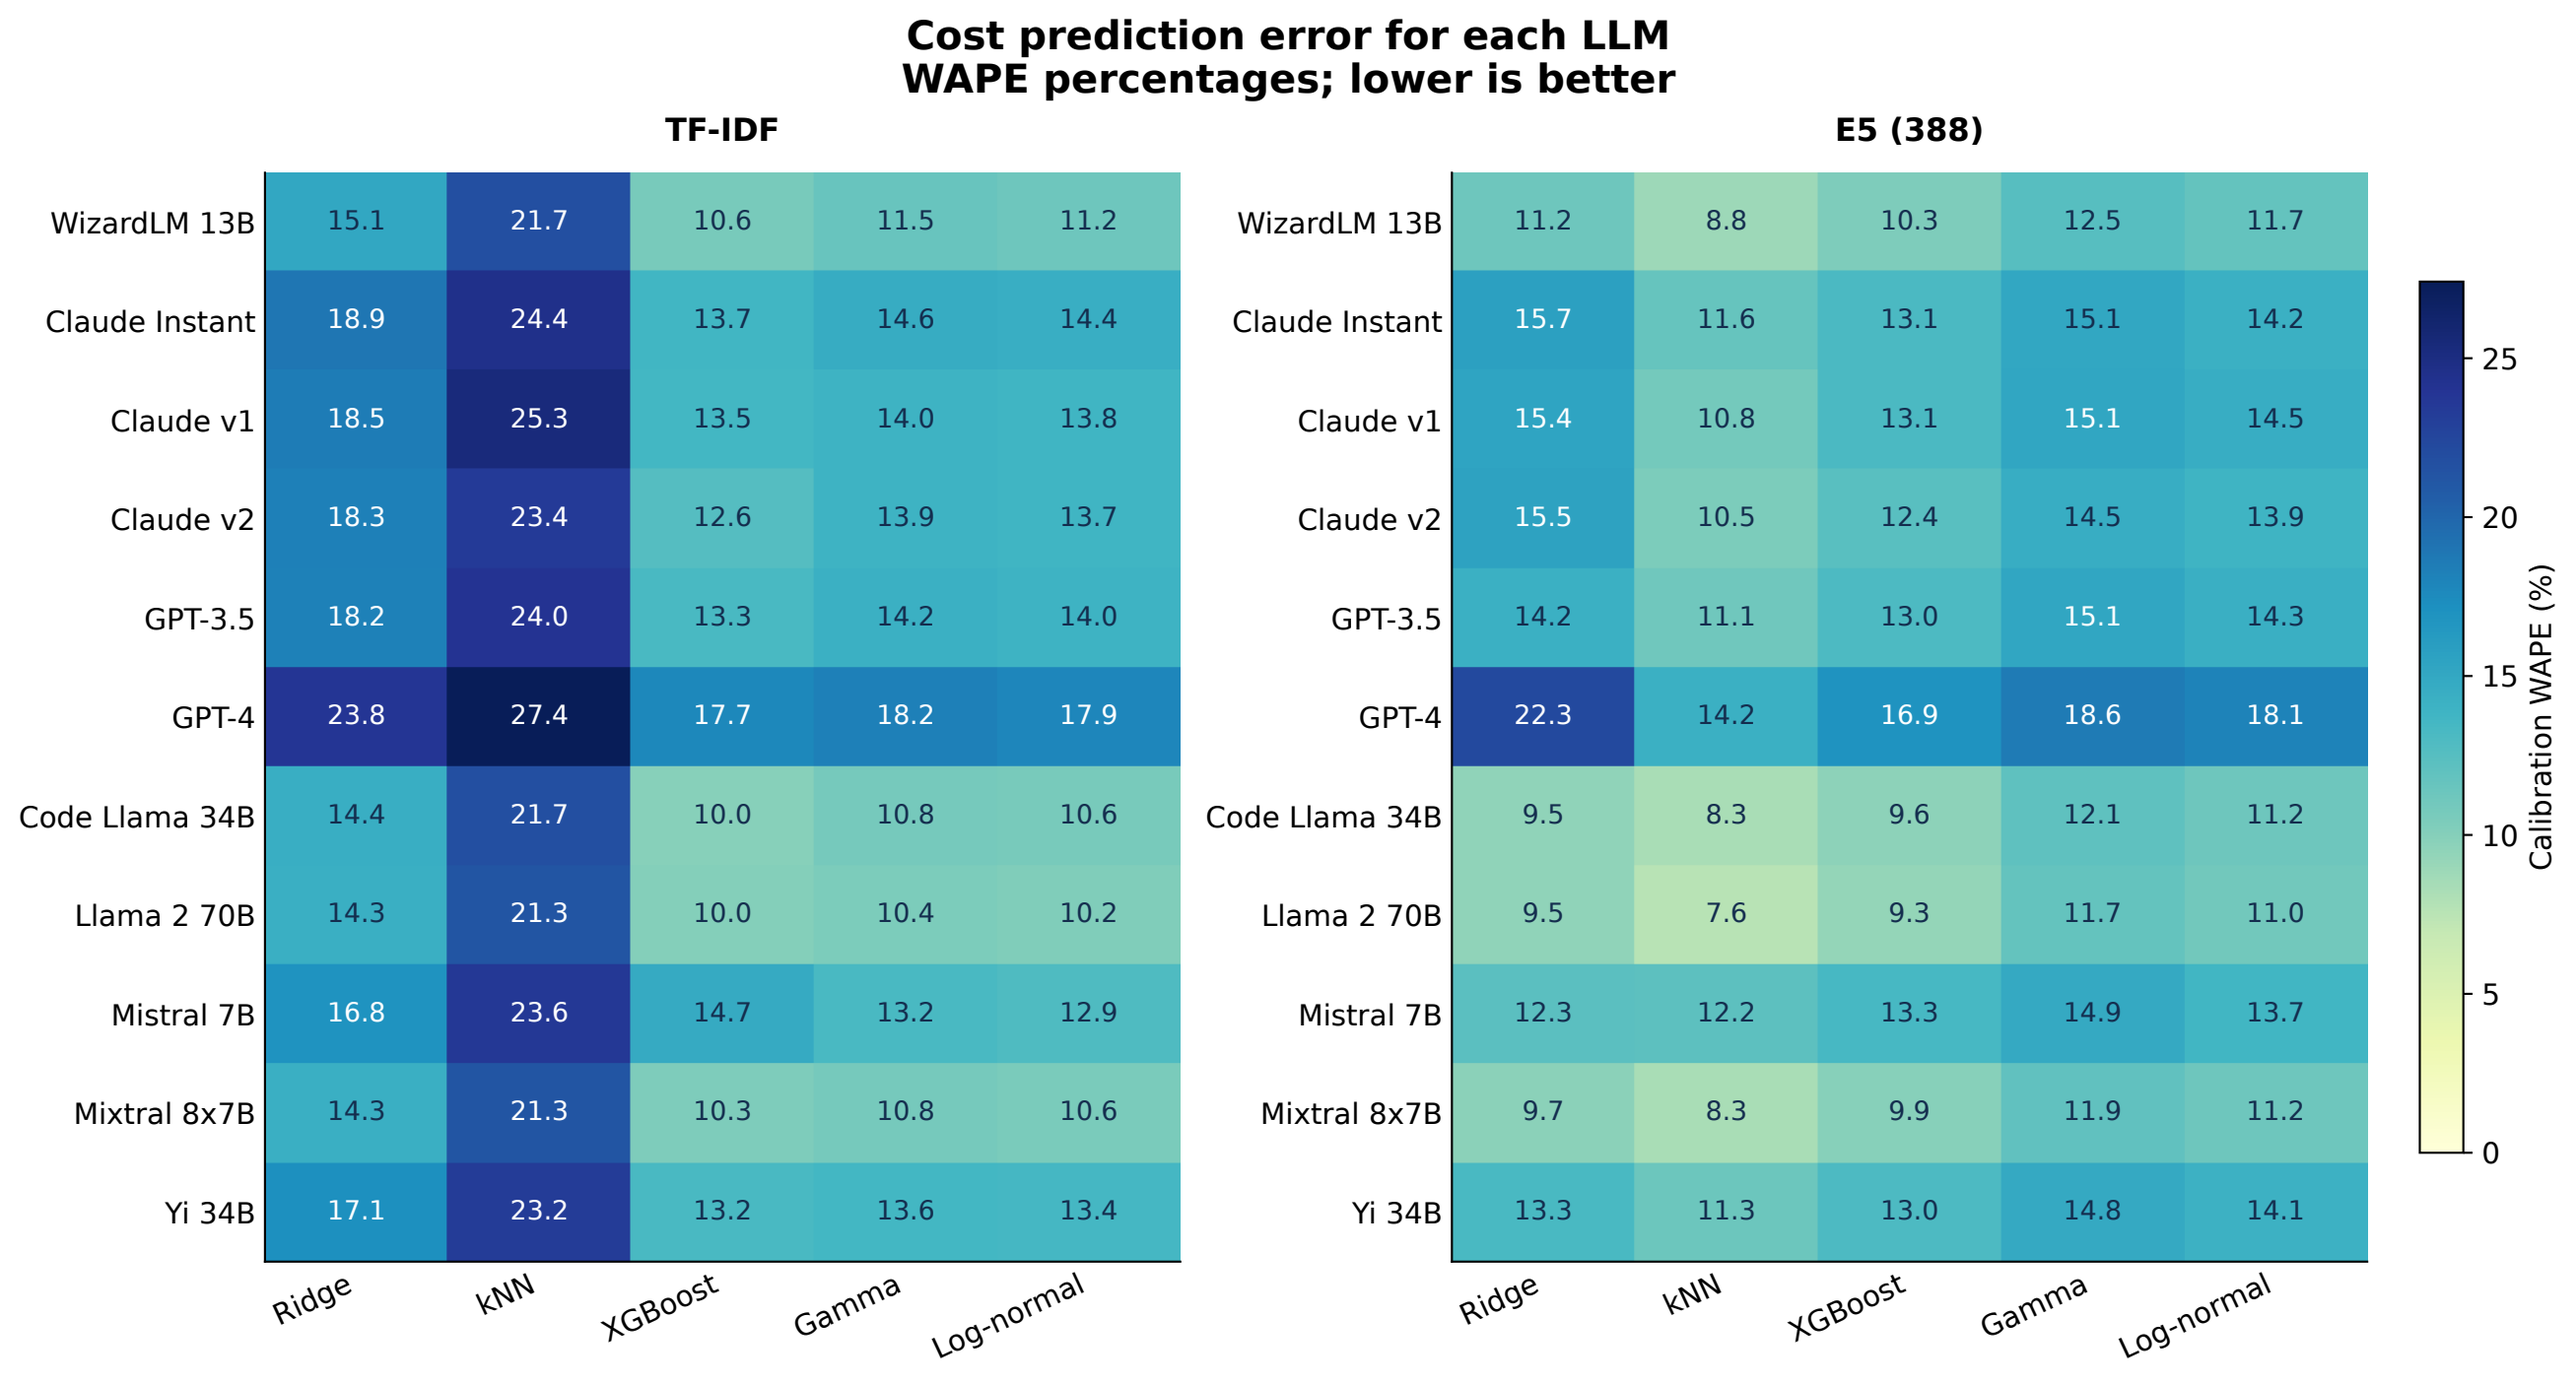
Cost prediction error for each LLM
WAPE percentages; lower is better
TF-IDF
E5 (388)
WizardLM 13B 15.1 21.7 10.6 11.5 11.2
Claude Instant 18.9 24.4 13.7 14.6 14.4
Claude v1 18.5 25.3 13.5 14.0 13.8
Claude v2 18.3 23.4 12.6 13.9 13.7
GPT-3.5 18.2 24.0 13.3 14.2 14.0
GPT-4 23.8 27.4 17.7 18.2 17.9
Code Llama 34B 14.4 21.7 10.0 10.8 10.6
Llama 2 70B 14.3 21.3 10.0 10.4 10.2
Mistral 7B 16.8 23.6 14.7 13.2 12.9
Mixtral 8x7B 14.3 21.3 10.3 10.8 10.6
Yi 34B 17.1 23.2 13.2 13.6 13.4
WizardLM 13B 11.2 8.8 10.3 12.5 11.7
Claude Instant 15.7 11.6 13.1 15.1 14.2
Claude v1 15.4 10.8 13.1 15.1 14.5
Claude v2 15.5 10.5 12.4 14.5 13.9
GPT-3.5 14.2 11.1 13.0 15.1 14.3
GPT-4 22.3 14.2 16.9 18.6 18.1
Code Llama 34B 9.5 8.3 9.6 12.1 11.2
Llama 2 70B 9.5 7.6 9.3 11.7 11.0
Mistral 7B 12.3 12.2 13.3 14.9 13.7
Mixtral 8x7B 9.7 8.3 9.9 11.9 11.2
Yi 34B 13.3 11.3 13.0 14.8 14.1
Ridge
kNN
XGBoost
Gamma
Log-normal
Calibration WAPE (%)
0
5
10
15
20
25